Aaditya Sangishetty

Software Engineer | Navi Mumbai, India | +91 7977672964
adityashetty35@gmail.com | linkedin.com/in/aaditya-sangishetty/ | github.com/adityashetty35 |
aadityasangishetty.vercel.app

SUMMARY

Software Engineer with 4+ years of experience designing and building scalable enterprise systems using Python and FastAPI. Experienced in multi-tenant architectures, distributed backend services, REST APIs, event-driven systems, real-time applications, and database-driven platforms serving 400+ enterprise clients and nearly 1 million requests per day. Hands-on experience with MySQL, MongoDB, Redis, NSQ, WebSockets, RBAC, and cloud-native infrastructure, along with building AI/LLM applications and MCP-based integrations for enterprise workflows.

CORE SKILLS

Languages: Python, Java, JavaScript/TypeScript

Backend Engineering: FastAPI, Flask, REST APIs, Microservices, API Design, Asynchronous Programming, OAuth, JWT

Distributed / Data: Redis, NSQ, RabbitMQ, PostgreSQL, MySQL, MongoDB, Idempotency, Retries, Batching, Caching, Event-Driven Systems, WebSockets

Frontend & Mobile: React, React Native, TypeScript, Vite

Applied AI: LLM Integration, MCP (Model Context Protocol), OCR, Document AI, RAG, LangChain, LlamaIndex, Azure OpenAI, Hugging Face

Cloud / Operations: Docker, GitHub Actions, Linux, Grafana, Prometheus

EXPERIENCE

Neural IT Pvt. Ltd. | Aug 2024 – Present
Python Developer | Navi Mumbai

- Led the redevelopment of a legacy PHP HRMS into a modern Python/FastAPI platform, designing a scalable multi-tenant backend architecture supporting 400+ enterprise clients.
- Contributed to the React web application and React Native mobile application for Android and iOS as part of the end-to-end redevelopment of the HRMS platform.
- Designed and implemented distributed backend services using Python/FastAPI, MySQL, MongoDB, Redis, and NSQ, supporting nearly 1 million requests per day across the HRMS platform.
- Optimized backend performance by implementing connection pooling, Redis caching, and asynchronous event processing, improving system scalability and reducing database overhead under concurrent workloads.
- Implemented real-time attendance and location tracking using FastAPI, Redis, Centrifugo WebSockets, and NSQ, enabling low-latency updates while asynchronously persisting tracking data.
- Built an AI/LLM platform with MCP server integration, implementing HRMS tools and capabilities that enable natural-language access to enterprise data and workflows.
- Designed and implemented authentication, authorization, and multi-tenant RBAC across the HRMS platform, supporting role-based access to modules, actions, and enterprise data.
- Containerized and deployed backend services using Docker and Kubernetes, implemented CI/CD automation and API management through Kong API Gateway, and established production monitoring and observability using Prometheus, Grafana, and Graylog.

Cerelabs | Jun 2022 – Aug 2024
Software Developer | Mumbai

- Maintained and enhanced an enterprise Document Masking Bot deployed for HDFC and Toyota Finance, supporting automated processing of confidential documents.
- Built machine learning models for Aadhaar card information extraction, achieving over 90% field extraction accuracy.
- Designed and implemented scalable backend services for high-volume document processing and enterprise automation workflows.
- Developed end-to-end backend components for Queryfic, a RAG-based enterprise knowledge platform, covering data ingestion, chunking, embeddings, vector search, retrieval, LLM integration, and conversational APIs for querying organizational data.

SELECTED PROJECTS

Hoshi HRMS Platform – Enterprise Multi-Tenant HRMS

- Re-engineered a legacy enterprise HRMS into a modern platform supporting 400+ clients, with tenant-isolated data, modular business services, and integrations across web and mobile applications.
- Implemented core HRMS domains, including payroll, attendance, leave, employee lifecycle, authentication, and access control, supporting configurable workflows across multiple clients.
- Designed and independently built reusable rule, authorization, reporting, and workflow engines to support configurable business rules, access control, reporting, and approval workflows across tenants.
- Built Hoshi as an extensible integration platform with SSO and MCP-based integrations, enabling external applications and AI clients to securely access HRMS data and capabilities. Integrated with analytics platforms such as Tableau for data retrieval and visualization.

Entrust Case Management System – Legal & Medical Case Management Platform

- Developed backend services for a microservices-based case management platform, implementing APIs for case, client, document, evidence, and workflow management across legal and medical business processes.
- Implemented Stripe payment gateway integrations, including automated payment flows and subscription-related billing workflows for the platform.
- Implemented asynchronous processing with NSQ and Redis-based caching to improve service responsiveness and efficiently handle background workloads across the platform.

Queryfic – Enterprise AI Knowledge Assistant

- Developed Queryfic, a RAG-based enterprise knowledge platform that enables users to query organizational data through natural-language conversations.
- Implemented the end-to-end RAG pipeline, covering data ingestion, document chunking, embeddings, vector search, retrieval, and context preparation for LLM-based responses.
- Integrated LLM-based response generation with retrieved organizational context, building conversational APIs that combine user queries with relevant knowledge to produce context-aware responses.

Document Masking Bot

- Developed an enterprise Document Masking Bot for automated identification, extraction, and masking of sensitive information from documents, deployed for financial-sector clients.
- Developed OCR and TensorFlow-based document processing pipelines to detect, extract, and automatically mask confidential information from unstructured documents.
- Built machine learning models for structured information extraction from identity documents, achieving over 90% field-level extraction accuracy.

EDUCATION

Bachelor of Science in Computer Science — 2019 – 2022
PVG College of Science & Technology — Mumbai
S K Somaiya Junior College — 2018 – 2019
Mumbai